

Efficient parallel C++

Final Project Presentation

Online Linear Equation Solving

Robert Andreas Fritsch robert.fritsch@student.kit.edu | February 27, 2023

Outline

1. Project Overview

2. Implementation Highlights

- Thread Pool
- Online Scheduler

3. Benchmarking

4. Conclusion and Outlook

Project Overview: Components

■ Solvers:

- Gaussian elimination (sequential)
- LU factorization (both sequential and parallel implementations)

■ Thread Pool:

- Constructed using multiple layers of queues

■ Scheduler:

- Trivial (processes all problems sequentially)
- Parallel (processes all problems in parallel)
- Mixed (selects strategy based on problem size and load)

■ Benchmarking and Tests:

- Comprehensive tests and timing routines (via `bench.hpp`) to assess performance.

Thread Pool

■ Task Queue Design:

1. A per-thread lock-free static (SPMC) queue for local subtasks.
2. A shared lock-free static (SPMC) queue for unassigned tasks.
3. A shared locked (MPMC) overflow queue that also allows threads to wait.

■ Overhead:

- Each worker logs binary timestamps for the start and end of tasks into a thread-local file.
The additional overhead is minimal and does not significantly affect benchmark results.

■ Implications:

- Under heavy load, subtasks are processed locally by each thread.
- The design permits threads to await or sleep when no tasks are available.

Online Scheduler: Strategies and Implications

■ Policy Variants:

- **Trivial Scheduler:** Enqueues every problem to be solved with the sequential LU solver.

When problems arrive with significant latency, many threads remain idle.


- **Parallel Scheduler:** Enqueues every problem to be solved with the parallel LU solver.

For small problems, the parallel LU solver is inefficient. In the long run, the trivial scheduler is more effective.



- **Mixed Scheduler:** Dynamically selects between the sequential and parallel LU solvers based on problem size and current load.

It achieves near-parallel performance for the problem cases considered. In theory, it should combine the advantages of both trivial and parallel scheduling, although this has not yet been conclusively demonstrated.


Benchmarking I






Benchmarking II



Benchmarking III



Benchmarking IV



Conclusion and Outlook

■ Summary:

- An efficient, low-overhead thread pool with logging enables the creation of detailed Gantt diagrams.
- The multi-layered queue design minimizes contention while providing robust fallback mechanisms.
- The scheduling strategies (Trivial, Parallel, Mixed) all perform well; however, in most cases, trivial scheduling proves to be the most efficient.

■ Future Work:

- Extend benchmarks across a broader range of hardware configurations.